

Notes while reading Reinforcement learning an introduction (Sutton/Barto)

Willem

January 2013

Contents

1	Introduction	1
2	Mutli-armed bandits	3
3	Finite Markov Decision Process	5
4	Dynamic Programming	7
4.1	Exercise 4.8	7
5	Monte Carlo Methods	9
5.1	Exercises	9
5.1.1	Exercise 5.1 page 94	9
5.1.2	Exercise 5.2 page 94	9
5.1.3	Exercise 5.4 page 99	9
5.1.4	Exercise 5.5 page 105	9
5.1.5	Exercise 5.6 page 108	10
5.1.6	Exercise 5.7 page 108	10
5.1.7	Exercise 5.8 page 108	10
5.1.8	Exercise 5.11 page 111	11
6	TD Prediction	13
6.1	Summary	13
6.1.1	TD prediction	13
6.1.2	TD Advantages	13
6.1.3	Optimality of TD(0)	14
6.1.4	SARSA	14
6.1.5	Q-Learning	14
6.1.6	Difference between SARSA and Q-Learning	14
6.1.7	Expected Sarsa	14
6.1.8	Double learning	15
6.2	Exercises	15

6.2.1	Exercise 6.1	15
6.2.2	Exercise 6.2	16
6.2.3	Exercise 6.3	16
6.2.4	Exercise 6.4	16
6.2.5	Exercise 6.5	16
6.2.6	Exercise 6.6	16
6.2.7	Exercise 6.7	17
6.2.8	Exercise 6.8	17
6.2.9	Exercise 6.11	17
6.2.10	Exercise 6.12	17
6.2.11	Exercise 6.13	17
6.2.12	Exercise 6.14	17
7	n-step bootstrapping	19
7.1	Summary	19
7.1.1	n-step TD prediction	19
7.1.2	n-step Sarsa	19
7.1.3	n-step Off-policy Learning	20
7.1.4	Per Decision Methods with Control Variates	20
7.1.5	Off-Policy Learning Without Importance Sampling: The n-step Tree Backup Algorithm	20
7.1.6	*A Unifying Algorithm: n-step Q(Sigma)	21
7.2	Exercises	21
7.2.1	Exercise 7.1	21
7.2.2	Excercise 7.3 page 145	22
7.2.3	Excercise 7.4 page 148	22

Chapter 1

Introduction

Chapter 2

Mutli-armed bandits

Chapter 3

Finite Markov Decision Process

Chapter 4

Dynamic Programming

4.1 Exercise 4.8

The reward is only obtained when the capital is above 99. When the capital is at 50, there is a 50% chance you can win the game. So this obviously is the optimal policy. When you reach 51: it would be rather odd to bet the entire capital, as you don't need to risk it all to reach 100. Bigger downside, but same upside. So the best course of action is to bet with 1, see if you can grow this above 50. If you lose it, you still have a 50% chance to win by betting it all.

Chapter 5

Monte Carlo Methods

5.1 Exercises

5.1.1 Exercise 5.1 page 94

The last 2 rows in the rear means you either have 21, or 20, which means the odd's are very good you will win. (hence high value function)

The last row on the left means the dealer has an ace, so it's at an advantage to get a higher score.

The front row's are higher on the upper diagram, as there is a usable ace. Which means that if you get a bad hit that put's you over 21. It can count as 1.

5.1.2 Exercise 5.2 page 94

As this is Markov process eg. The cards drawn are not exhaustible. The odds of winning on the second time your in the same state is just as good as the first time.

5.1.3 Exercise 5.4 page 99

The "Append G to Returns (S_t, A_t) would be replaced by increasing a count and added it as running average to some table.

5.1.4 Exercise 5.5 page 105

question: Consider an MDP with a single Non-terminal state and a single action that transitions back to the nonterminal state with probability p and transitions to the terminal state with probability

$p-1$. Let the reward be 1 on all transitions, and let $\gamma = 1$. Suppose you observe one episode that lasts 10 steps, with a return of 10. What are the first-visit and every visit estimators of the value of the non-terminal state.

10 Steps means 9 towards the non-terminal, and one towards the terminal. The rewards are all-way's the same so the final cost=10.

If $\gamma = 1$ then $G = G + \gamma R_{k+1}$ in every iteration.

In case of all visit the complete horizon counts 10 times in the non-terminal state, as the 10th time we leave the non-terminal state for good and enter the terminal state. $(1+2+3+4+5+6+7+8+9+10)/10 = 55/10 = 5.5$ So the value is 5.

In case of the first-visit, we only count the first visit which has a reward of 1.

5.1.5 Exercise 5.6 page 108

question: What is the equation analogous to (5.6) for action values $Q(s, a)$ instead of state values $V(s)$, again given returns generated using b?

$Q(s, a)$ is similar to $V(s)$, it takes the $V(s)$ given a certain step was taken first.

$$Q(s, a) = \frac{\sum_{t \in J(s, a)} \rho_{t+1:T(t)-1} G_t}{\sum_{t \in J(s, a)} \rho_{t+1:T(t)-1}} \quad (5.1)$$

5.1.6 Exercise 5.7 page 108

question: In learning curves such as those shown in Figure 5.3 error generally decreases with training as indeed happened for the ordinary importance-sampling method. But for the weighted importance-sampling method error first increased and then decreased. Why do you think this happened

If there are but a few samples, the bias will be the dominating error. And it will increase as more and more samples are added. Until there are so many samples, it starts to disappear.

5.1.7 Exercise 5.8 page 108

question: The results with Example 5.5 and shown in Figure 5.4 used a first-visit MC method. Suppose that instead an every-visit MC method was used on the same problem. Would the variance

of the estimator still be infinite? Why or why not? A first Visit MC has less terms than a every Visit MC. All terms have a positive value, so it would also go to infinite.

5.1.8 Exercise 5.11 page 111

If the target policy is a greedy deterministic policy, and the loop is broken off if $\pi(S_t) \neq A_t$. Then $\pi(A_t|S_t) = 1$ by definition.

Chapter 6

TD Prediction

6.1 Summary

6.1.1 TD prediction

The basic formula for monte carlo prediction is $V(S_t) = V(S_t) + \alpha[G_t - V(S_t)]$. G_t is the final result, this means that the update only can happen at the end of the simulation. By replacing G_t with $R_{t+1} + \gamma V(S_{t+1})$ we get the TD method $V(S_t) := V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$.

The update of the TD method is called the **TD error** $\delta = G_t - R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$. An equivalent entity exists with monte carlo methods, and is called the monte-carlo error. The monte carlo error can be written as a sum of TD errors, illustrated by equation 6.1. (proof on page 121)

$$G_t - V(S_t) = \sum_{k=t}^{T-1} \gamma^{k-t} \delta_k \quad (6.1)$$

6.1.2 TD Advantages

1. No model of the behavior is required
2. Naturally online/incremental algorithm (useful with long episodes)
3. Learns from experimental choices (monte carlo need to discard them)
4. In practice faster than monte carlo methods

6.1.3 Optimality of TD(0)

When using batch learning, as in only changing the value function everytime a whole batch is processed. TD(0) and Monte Carlo do not converge to the same solution. Monte Carlo methods find the solution that minimized the error on the dataset. TD(0) finds the parameters that most likely would cause a Markov process to result in the dataset. This is called the certainty-equivalence estimate.

6.1.4 SARSA

SARSA stands for $S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}$. It uses a policy to generate A_t and A_{t+1} . Updates the Q value, applies A_{t+1} and then finds the next input A_{t+2} .

$$Q(S_t, A_t) := Q(S_t, A_t) + \alpha[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)] \quad (6.2)$$

6.1.5 Q-Learning

Q-learning acts greedily in when predicting, but acts according to its policy when finding an input to apply to the system. So in contrast to SARSA it won't reuse A_{t+1} it generated when predicting.

$$Q(S_t, A_t) := Q(S_t, A_t) + \alpha[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)] \quad (6.3)$$

6.1.6 Difference between SARSA and Q-Learning

SARSA will act a bit more careful, as its prediction is not greedy. And it takes into account that the next action might not be the best one. Q-Learning will take the more risky route, as it uses the best (according to $Q(S, A)$) possible action in its prediction.

6.1.7 Expected Sarsa

Expected SARSA uses the expected value of all possible actions A_{t+1} given the policy. Then it uses a greedy policy to act, just like with Q-learning. Expected Sarsa will work with $\alpha = 1$, which would not work very well with classical SARSA. This makes the short term behavior much better. But is more computationally expensive.

$$Q(S_t, A_t) := Q(S_t, A_t) + \alpha[R_{t+1} + \gamma \mathbb{E}[Q(A_{t+1}, S_{t+1})|S_{t+1}] - Q(S_t, A_t)] \quad (6.4)$$

6.1.8 Double learning

Equation 6.3 uses an argmax to estimate the value of Q . If one of these estimates is over-estimated, it will result in bad behavior(bias). Double learning reduces the odds of this happening by using two $Q(A, S)$ estimates. One to find the maximum action, and one to estimate it's value.(equation 6.5) It's less like that the overestimate will happen this way.

$$A = Q_2(\arg \max_a Q_1(S, a)) \quad (6.5)$$

It's good practice to swap Q_1 and Q_2 in equation 6.5 constantly. For example at random with odds 50/50.

6.2 Exercises

6.2.1 Exercise 6.1

$$V_{t+1}(s_t) = \alpha[R_{t+1} + \gamma V_t(s_{t+1}) - V_t(s_t)] + V_t(s_t) \quad (6.6)$$

The difference between the value function at time t and $t+1$ is defined by equation 6.6.

The equality $G_t = R_{t+1} + \gamma G_{t+1}$ still holds. However the monte carlo error is slightly different in every iteration. $G_t - V_t(s_t)$ becomes $G_{t+1} - V_{t+1}(s_{t+1})$ in the next iteration. As the value function now changes at iteration t , with a difference of $d_t = \alpha[R_{t+1} + \gamma V_t(s_{t+1}) - V_t(s_t)]$.

$$G_{t+1} - V_t(S_{t+1}) = G_{t+1} - V_{t+1}(S_{t+1}) - d_{t+1} \quad (6.7)$$

$$error = - \sum_{k=t+1}^{T-1} \gamma^{k-t} d_{k-1} \quad (6.8)$$

In conclusion the different factor is equation 6.8.

6.2.2 Exercise 6.2

If (as explained in the example of the hint) a part of the statespace is already well estimated. Then the TD prediction will be very good as you enter those states and if your path ends on one of those states. So you only have lesser predictions while in an unexplored part.

The Monte Carlo approach would still need to evaluate through the already well estimated part. Which is rather slow.

6.2.3 Exercise 6.3

The change on a value function is defined by: $\alpha[R_{t+1} + \gamma V_t(s_{t+1}) - V_t(s_t)] = 0.1[0 + 0 - 0.5] = -0.05$ if $V_t(s_{t+1}) = 0$ so it ends on the left terminal state. And $\alpha = 0.1$ and $V_t(A) = 0.5$.

6.2.4 Exercise 6.4

The TD algo is over-fitting when $\alpha > 0.05$ we could try to make it a bit smaller. But at $\alpha = 0.05$ it seems to flatten out nicely, so I would not expect better results.

A similar story with the MC method, this time at $\alpha 0.02$ we get a nice flat tail. It's not as clear as with the TD method, but that's due the larger variance on the MC method.

So no, I would not expect any changes in results if more samples were ran with different values for α .

6.2.5 Exercise 6.5

Overfitting, the step is too large so TD cannot find the optimal values. But keeps over/under estimating every time it runs through an episode.

6.2.6 Exercise 6.6

You setup the bellman optimality equation, and the pick a method to solve it. As this is a rather simple example, you could just manually solve the equation.

$$\begin{aligned}
V(A) &= 0.5V(B) \\
V(B) &= 0.5V(A) + 0.5V(C) \\
V(C) &= 0.5V(B) + 0.5V(D) \\
V(D) &= 0.5V(C) + 0.5V(E) \\
V(E) &= 0.5V(D) + 0.5
\end{aligned}
\tag{6.9}$$

This seems like the simplest way to do it, as it's small.

6.2.7 Exercise 6.7

The normal on-policy TD(0) update looks like $V(s_t) = V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$. I would expect that $\alpha = \frac{\rho}{\sum_t \rho^t}$ as it becomes a weighted average due to the importance sampling.

6.2.8 Exercise 6.8

todo, not hard, but a bit of bookkeeping to be done.

6.2.9 Exercise 6.11

In Q-learning the actions that are applied to the system are learning through a ϵ -greedy policy (behavior policy) are not used for the prediction(Q). This is by definition an off-policy control.

6.2.10 Exercise 6.12

It would be nearly the same, SARSA selects the next action before updating Q and Q-learning selects it after. So the update of Q might make a difference in some cases.

6.2.11 Exercise 6.13

todo

6.2.12 Exercise 6.14

todo

Chapter 7

n-step bootstrapping

7.1 Summary

7.1.1 n-step TD prediction

Monte Carlo updates the estimate $V(S)$ using the complete return (equation 7.1). TD(0) only watches one step ahead, a compromised is too take an n-step prediction window as illustrated by equation 7.2. The state learning algorithm then becomes equation 7.3.

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t-1} R_T \quad (7.1)$$

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n}) \quad (7.2)$$

$$V_{t+n}(S_t) = V_{t+n-1}(S_t) + \alpha [G_{t:t+n} - V_{t+n-1}(S_t)] \quad (7.3)$$

7.1.2 n-step Sarsa

The previous Sarsa is often called Sarsa(0), the generalized version is called n-step Sarsa. The return value can be estimated by equation 7.4. The update rule for Q then becomes equation 7.5.

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} \dots \gamma^{n-1} R_{t+n-1} + \gamma^n Q_{t+n-1}(S_{t+n}, A_{t+n}) \quad (7.4)$$

$$Q_{t+n}(S_t, A_t) := Q_{t+n-1}(S_t, A_t) + \alpha [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)] \quad (7.5)$$

The same logic can be applied to expected SARSA. The value function can be defined as equation 7.6. Using equation 7.6 the n-step expected sarsa

return value is defined by equation 7.7. The update is still equation 7.5 but using the new $G_{t:t+n}$ from equation 7.7.

$$V_t(S_t) = \sum_a P(a)Q(S_t, a) \quad (7.6)$$

$$G_{t:t+n} = R_t + \gamma R_{t+1} + \dots \gamma^{n-1} R_{t+n-1} + V_t(S_{t+n}, A_{t+n}) \quad (7.7)$$

7.1.3 n-step Off-policy Learning

$$Q_{t+n}(S_t, A_t) := Q_{t+n-1}(S_t, A_t) + \rho_{t:t+n} \alpha [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)] \quad (7.8)$$

$$V_{t+n}(S_t) = V_{t+n-1}(S_t) + \rho_{t:t+n-1} \alpha [G_{t:t+n} - V_{t+n-1}(S_t)] \quad (7.9)$$

$$\rho_{t:h} = \prod_{k=t}^{\min(h, T-1)} \frac{\pi(A_k | S_k)}{b(A_k | S_k)} \quad (7.10)$$

7.1.4 Per Decision Methods with Control Variates

The previous sections are inefficient implementations of the n-step algorithm. The n-step return can be written recursively as $G_{t:h} = R_{t+1} + \gamma G_{t+1:h}$, with $G_{h:h} = V_{h-1}(S_h)$. The importance sampling weighting is still $\rho_t = \frac{\pi(S_t, A_t)}{b(S_t, A_t)}$. Using this recursive definition we can define the return as equation 7.11.

$$G_{t:h} = \rho_t(R_{t+1} + \gamma G_{t+1:h}) + (1 - \rho)V_{h-1}(S_t) \quad (7.11)$$

The term $(1 - \rho)V_{h-1}(S_t)$ in equation 7.11 is called the control variate. It has an expected value of one because $\mathbb{E}[\rho] = 1$ so $\mathbb{E}[1 - \rho] = 0$.

The return state of a n-step off policy with control variate is defined by equation 7.12. The recursion ends with $G_{h:h} = Q_{h-1}(S_h, A_h)$

$$G_{t:h} = R_{t+1} + \gamma[\rho_{t+1}G_{t+1:h} + V_{h-1}(S_{t+1}) - \rho_{t+1}Q_{h-1}(S_{t+1}, A_{t+1})] \quad (7.12)$$

7.1.5 Off-Policy Learning Without Importance Sampling: The n-step Tree Backup Algorithm

Instead of only using the value estimations of the actual path taken in n-steps. We can generalize the expected SARSA algorithm, from equation 7.13 to equation 7.14.

$$G_{t:t+1} = R_{t+1} + \gamma \sum_a \pi(a|S_{t+1}) Q_t(S_{t+1}, a) \quad (7.13)$$

$$\begin{aligned} G_{t:t+n} &= R_{t+1} + \gamma \sum_{a \neq A_{t+1}} \pi(a|S_{t+1}) Q_{t+n-1}(S_{t+1}, a) + \gamma \pi(A_{t+1}|S_{t+1}) G_{t+1:t+n} \\ &= R_{t+1} + \gamma \pi(A_{t+1}|S_{t+1}) [G_{t+1:h} - Q_{h-1}(S_{t+1}, A_{t+1})] + \gamma V_{h-1}(S_{t+1}) \end{aligned} \quad (7.14)$$

Adding the control variate to Equation 7.14 gives equation 7.15.

$$G_{t:t+n} = R_{t+1} + \gamma \pi(\sigma_{t+1} \rho_{t+1} + (1 - \sigma_{t+1}) \pi(A_{t+1}|S_{t+1})) [G_{t+1:h} - Q_{h-1}(S_{t+1}, A_{t+1})] + \gamma V_{h-1}(S_{t+1}) \quad (7.15)$$

7.1.6 *A Unifying Algorithm: n-step Q(Sigma)

7.2 Exercises

7.2.1 Exercise 7.1

The Monte carlo error can be written as a sum of TD errors, with TD(0) this becomes:

$$\begin{aligned} G_t &= R_{t+1} + \gamma G_{t+1} \\ \delta_t &= R_{t+1} + \gamma V(S_{t+1}) - V(S_t) \\ G_t - V(S) &= R_{t+1} + \gamma G_{t+1} - V(S_t) = \sum_{k=t}^{T-1} \gamma^{k-t} \delta_k \end{aligned}$$

With an n step we get:

$$\begin{aligned} G_t &= R_{t+1} + \gamma G_{t+1} \\ \delta_t &= \sum_{k=1}^n \gamma^{k-1} R_{t+k} + \gamma^n V(S_{t+n}) - V(S_t) \end{aligned}$$

By putting them together we get:

$$\begin{aligned} &G_t - V(S_t) \\ &= R_{t+1} + \gamma G_{t+1} - V(S_t) \\ &= R_{t+1} + \gamma G_{t+1} - V(S_t) \\ &+ \sum_{k=2}^n \gamma^{k-1} R_{t+k} - \sum_{k=2}^n \gamma^{k-1} R_{t+k} \\ &+ \gamma^n V(S_{t+n}) - \gamma^n V(S_{t+n}) \\ &= \delta_t + \gamma(G_{t+1} - \gamma^{n-1} V(S_{t+n})) - \sum_{k=2}^n \gamma^{k-1} R_{t+k} \end{aligned}$$

I don't see how to continue from here.

7.2.2 Exercise 7.3 page 145

question: Why do you think a larger random walk task (19 states instead of 5) was in the examples of this chapter? Would a smaller walk have shifted the advantage to a different value of n ? How about the change in left-side outcome from 0 to -1 made in the larger walk? Do you think that made any difference in the best value of n ?

If a random walk with length of 5 was used the results would not be best around an n of 4 or 8. As you need a large enough episode to learn from a step of 4 or 8 samples.

When the walk is smaller, smaller n 's will give better results than current results.

If the left side is negative, the first time the walk will go into it. That negative result will propagate n -steps, instead of 1 with $TD(0)$.

7.2.3 Exercise 7.4 page 148

Similar to exercise 7.1, still TODO, first finish 7.1.